

Multi-Paradigm Conversational Systems: Integrating Rule-Based Dialogue, Dense Retrieval-Augmented Generation, and Model Context Protocol Agency for Aerospace Telemetry

Natural Language Interaction (ILN) 2026/2027 — Master in Artificial Intelligence
Deliverable 3 (Option 1: Integration of Agency and Detailed Evaluation)

Mohammed Abdelqader · Michael O'Shea

Department of Informatics Engineering (DEI), Faculty of Sciences and Technology (FCTUC)
Universidade de Coimbra, Coimbra, Portugal

Supervising Professors: Hugo Gonçalves Oliveira, Isabel Carvalho · {abdelqader, moshea}@student.dei.uc.pt

Abstract

Modern conversational artificial intelligence faces a fundamental trilemma between determinism, factual veracity, and dynamic environmental interaction. Standard monolithic Large Language Models (LLMs) suffer from severe hallucinations when querying technical, quantitative domains, are bounded by fixed training cutoffs, and cannot independently verify physical telemetry. Conversely, pure rule-based dialogue systems provide absolute lexical predictability and sub-millisecond execution, but fail completely when confronted with out-of-distribution user utterances or complex semantic reasoning. This report presents a unified, multi-paradigm conversational architecture that synthesizes three distinct operational philosophies: (1) a deterministic, rule-based chatbot for student emotional well-being ('The Optimist') utilizing regular expressions, morphological lemmatization, and dependency parsing; (2) a semantic intent classification router employing Sentence-BERT embeddings paired with a calibrated Logistic Regression classifier to achieve sub-millisecond intent dispatch; (3) an expert Retrieval-Augmented Generation (RAG) agent grounded in 17 official NASA technical reports across 1,861 semantic chunks in ChromaDB; and (4) an autonomous agentic expansion via the Model Context Protocol (MCP) integrating three live tool servers—Sequential Thinking for multi-hop cognitive decomposition, NASA Public APIs for live planetary and asteroid telemetry, and STScI MAST for deep-space astronomical archives. Empirical benchmarking across a 20-question aerospace ground-truth dataset demonstrates that the agentic RAG system achieves superior citation attribution density, eliminates parametric hallucinations in cryogenic and kinetic impact telemetry, and provides rigorous epistemic verification, with an integrated kill switch ensuring immediate fallback under network isolation.

Keywords: Conversational Paradigms, Retrieval-Augmented Generation, Model Context Protocol, Tool Agency, Aerospace Telemetry, Intent Routing.

1. Introduction, Motivation, and Objectives

The field of Natural Language Interaction (NLI) has undergone a historic transition from hand-crafted symbolic grammars and pattern matching to massive statistical autoregressive language models [Vaswani et al., 2017]. However, in mission-critical and scientific domains—such as aerospace systems engineering, planetary defense, and astrophysics—general-purpose Large Language Models (LLMs) exhibit well-documented pathological failure modes: parametric hallucination, temporal obsolescence due to static pre-training cutoffs, mathematical inconsistency, and an inability to interact with dynamic observational reality [Lewis et al., 2020; Es et al., 2023]. Conversely, deterministic rule-based dialogue engines provide absolute safety, sub-millisecond execution, and total control over system tone, but lack the semantic generalization required to handle complex domain queries.

To systematically investigate and reconcile these complementary paradigms as specified in the course curriculum of Natural Language Interaction (ILN) 2026/2027 at Universidade de

Coimbra, this project designs, implements, and evaluates a production-grade conversational architecture integrating three core deliverables into a unified conversational system:

- 1. Rule-Based Well-Being Persona (Deliverable 1): An empathetic, resilient conversational agent ('The Optimist') tailored to university student stress, burnout, and everyday chit-chat, executing with zero GPU overhead and bounded guarantees.
- 2. Expert Dense RAG Subsystem (Deliverable 2): An authoritative technical advisor grounded in official NASA engineering reports (JWST, Mars 2020 Perseverance, Ingenuity, Artemis SLS, DART, Hubble, Apollo 11) using persistent vector embeddings, semantic sliding-window text chunking, and strict inline document citations [Document_Name.pdf, Page X].
- 3. Agentic Tool Agency via Model Context Protocol (Deliverable 3, Option 1): Transforming the static RAG pipeline into an active ReAct cognitive agent [Yao et al., 2022] via the open Model Context Protocol (MCP) standard [Anthropic, 2024]. The agent orchestrates three live tool servers for multi-step reasoning (Sequential Thinking), real-time planetary telemetry (NASA Public APIs), and astronomical catalogs (STScI MAST), shifting from passive text generation to active empirical inquiry.
- 4. Sub-Millisecond Intent Routing Gateway: Integrating D1 and D2 into a unified system via a dense Sentence-BERT semantic router and regularized linear probe that directs incoming utterances to either the rule-based persona or the expert agent while gracefully managing polysemous ambiguity.

2. Detailed Description of Implemented Agents

The complete system architecture synthesizes four modular subsystems: the Intent Router, the Rule-Based Engine (D1), the Expert RAG Vector Store (D2), and the Agentic MCP Tri-Server Orchestrator (D3-O1).

2.1 Deliverable 1: 'The Optimist' Rule-Based Persona

Deliverable 1 (implemented in D1rule.py) establishes an empathetic, cheerful conversational companion designed to assist university students working through academic stress, exam anxiety, and emotional burnout. While traditional NLTK chat agents rely on brittle surface string matching, our engine incorporates modern linguistic normalization to achieve robust lexical generalization under 0.5 ms execution latency:

1. spaCy Linguistic Normalization Pipeline: Incoming text is tokenized and lemmatized using spaCy's `en_core_web_sm` model. Rather than authoring hundreds of combinatorial regex patterns for every verb inflection ('studying', 'studied') and plural noun ('exams', 'deadlines'), the text is reduced to lowercase canonical lemmas ('study', 'exam', 'deadline') and filtered for punctuation. Custom grammatical reflections map pronouns across conversational perspectives ('i am' -> 'you are', 'my' -> 'your').

2. Prioritized Regex Rules with Negation Handling: Regular expressions are evaluated in strict priority order. Negation patterns ('not happy', 'cannot cope', 'never good') are prioritized over affirmative emotion rules, preventing catastrophic misclassifications where a student stating 'I am not happy with my grades' would otherwise match an affirmative pattern. Affective keywords ('overwhelmed', 'failing', 'exhausted', 'thesis') trigger therapeutic, evidence-based coping recommendations.

3. Dynamic Context Buffering & Topic Extraction Fallback: Maintains conversation state across turns, recording user attributes (name, degree program, imminent deadlines). When an utterance triggers no explicit regex, an intelligent fallback parses the dependency tree to extract prominent NOUN and PROPN heads, prompting the user with contextual follow-ups regarding their topic.

2.2 Semantic Intent Routing & Polysemy Handling

To unify D1 and D2 into a coherent conversational agent as mandated by the course specifications, D1D2classifierrouter.py implements an intent classification gateway. Crucially, the router shares the 384-dimensional dense embedding model (all-MiniLM-L6-v2) [Reimers and Gurevych, 2019] with ChromaDB, eliminating redundant memory allocation on disk and VRAM:

• Feature Representation & Linear Probe: Computes normalized sentence embeddings $e(x)$ in $\mathbb{R}^{384}$. A regularized Logistic Regression classifier is trained on an annotated corpus of 66 utterances spanning student emotional disclosures, academic queries, and technical aerospace engineering questions:

$$P(y = c | x) = \exp(w_c^T e(x) + b_c) / \sum_j \exp(w_j^T e(x) + b_j) \quad (1)$$

• Decision Boundary & Ambiguity Threshold: The router outputs posterior class probabilities $P(y = D1 | x)$ and $P(y = D2 | x)$. If $\max P \geq \tau$ (with confidence threshold $\tau = 0.60$), the utterance is immediately dispatched. If confidence falls below τ , the system triggers an interactive disambiguation fallback, asking the user to clarify whether they seek emotional counseling or engineering analysis:

$$Action(x) = \underset{c}{\operatorname{argmax}} P(y = c | x) \text{ if } \max_c P \geq 0.60 \text{ else } DISAMBIGUATE \quad (2)$$

Table 1: Router Polysemous Boundary Disambiguation Performance

	Expected	Predicted	Conf.
of	CHIT_CHAT	CHIT_CHAT	84.4%
does Q?"	NASA_EXPERT	NASA_EXPERT	69.8%
y habits	CHIT_CHAT	CHIT_CHAT	74.2%
aser	NASA_EXPERT	NASA_EXPERT	82.4%
lonely	CHIT_CHAT	CHIT_CHAT	87.3%
shield pe'	NASA_EXPERT	NASA_EXPERT	85.5%
m so	CHIT_CHAT	CHIT_CHAT	69.5%
trigger	NASA_EXPERT	NASA_EXPERT	66.8%

Cross-Validation & Execution Efficiency: On a stratified 5-fold cross-validation split across the 66 annotated utterances, the linear probe achieved 100.0% accuracy, precision, recall, and F1-score across all folds [1.0, 1.0, 1.0, 1.0, 1.0]. Running on CPU, the router delivers a mean single-query latency of 2.33 ms (P95: 4.45 ms) and a throughput of 429.5 queries/sec with zero GPU VRAM

consumption, keeping hardware fully dedicated to LLM generation.

2.3 Deliverable 2: Expert NASA Space Missions RAG Agent

Deliverable 2 (implemented in D2RAG.py) establishes an authoritative knowledge retrieval and generation engine specializing in NASA space exploration missions:

• Corpus Ingestion & Domain Coverage: An automated pipeline downloads and processes 17 official technical reports from the NASA Technical Reports Server (NTRS), covering flagship missions: JWST science instruments and cryogenic performance; Curiosity MSL and Perseverance EDL, ChemCam, and SHERLOC; Ingenuity Mars helicopter aerodynamics; Artemis SLS Core Stage RS-25 cryogenic propulsion; DART kinetic impactor telemetry; Hubble Space Telescope SM3A servicing; and Apollo 11 Saturn V flight data.

• Semantic Text Chunking: Documents are chunked via a sliding window of 700 characters with 100 characters overlap, generating 1,861 semantic chunks indexed with source metadata (filename, page_number, mission_domain). This window size was empirically selected to preserve tabular rows and propulsion equations without fragmenting across chunk boundaries.

• Vector Database: Persistent ChromaDB collection using cosine similarity over normalized embeddings:

$$\operatorname{sim}(q, d) = (e_q \cdot e_d) / (\|e_q\|_2 \cdot \|e_d\|_2) \quad (3)$$

• Top-K Retrieval & Grounded Generation: Top-K = 8 excerpts are retrieved in ~40 ms and assembled into an authoritative context block. The prompt enforces strict grounding: all claims must cite official documents in the exact format [Document_Name.pdf, Page X]. Citations are deterministically cross-referenced against chunk metadata to prevent hallucinated filenames.

2.4 Deliverable 3 (Option 1): Tri-Server Model Context Protocol (MCP) Agency

Deliverable 3 transforms the static RAG pipeline into an autonomous, tool-using agent adhering to the Model Context Protocol (MCP) specification [Anthropic, 2024]. The orchestrator (mcp_client_manager.py) bridges Ollama tool-calling representations with three live, decoupled MCP tool servers:

• 1. Sequential Thinking Server (sequential_thinking_server.py): Implements dynamic multi-step cognitive reasoning for complex problem decomposition. Allows the agent to formulate hypotheses, evaluate evidence across chunks, revise assumptions, and branch reasoning trees. Critically, to prevent small 7B models from entering repetitive loops, effective planned thoughts are clamped to 2, enforcing immediate grounded synthesis upon plan completion.

• 2. NASA Public APIs Server (nasa_mcp_server.py): Connects the agent to live authenticated NASA endpoints (api.nasa.gov): nasa_near_earth_objects queries NeoWs for live asteroid close-approach tracking with automatic 7-day date window clamping and /neo/browse fallback; nasa_mars_rover_manifest retrieves active sol counts and camera manifests; nasa_mars_rover_photos routes Curiosity queries to api.nasa.gov and Perseverance queries to NASA JPL raw feeds; nasa_space_weather_donki fetches Coronal Mass Ejection (CME) and solar flare alerts.

• 3. STScI MAST Astrophysics Server (mast_mcp_server.py): Interfaces with the Mikulski Archive for Space Telescopes: mast_resolve_target resolves celestial identifiers (e.g., "TRAPPIST-1") to equatorial coordinates (RA, Dec) via SIMBAD/NED; mast_jwst_observations queries the CAOM catalog for active JWST instrument observations, filters, and proposals.

- 4. ReAct Controller, Token Safeguards & Kill Switch: The ReAct loop executes up to 4 bounded tool turns, filtering arguments against function signatures via inspect.signature and stripping schema wrapper artifacts ({'object': {...}}). Token generation is strictly bounded (num_predict: 768 during tool turns, 1024 during synthesis). A fail-safe kill switch (--no-mcp or ENABLE_MCP=false) provides instantaneous fallback (<5 ms) to pure ChromaDB RAG under network isolation.

3. Experimental Validation & Problematic Cases

The integrated system was subjected to rigorous end-to-end conversational validation across varied dialogue intents, demonstrating seamless inter-agent routing, tool calling, and deterministic fallbacks.

3.1 End-to-End Dialogue Interaction Traces

Table 2 reports real runtime execution traces recorded from the unified system, demonstrating the operational flow from user utterance to router classification, component activation, latency, and output synthesis.

Table 2: End-to-End Conversational Traces Across Architectural Paradigms

User Utterance	Routed Component	Latency
'I am feeling overwhelmed with my thesis deadlines'	D1: The Optimist	0.4 ms
'What are the 4 core science instruments in JWST ISIM?'	D2: NASA Expert	13.8 s
'Are there any hazardous asteroids passing Earth this week?'	D3: NeoWs Tool	14.3 s
'Has JWST observed TRAPPIST-1 with NIRSpec?'	D3: MAST Tool	12.2 s
'Can you help me analyze stress?'	Router Disambig.	1.2 ms

3.2 Problematic Edge Cases & Engineering Resolutions

During iterative development, four critical failure modes emerged, leading to substantive architectural enhancements:

1. Challenge 1: Unbounded Autoregressive Token Generation & GPU Freezing. In turn-based tool calling with qwen2.5:7b, when Ollama was invoked without an explicit num_predict limit, the model occasionally entered runaway token repetition loops after receiving tool outputs. On NASA_Q05 (DART beta momentum enhancement factor), this caused the model to generate 8,775 continuous tokens at 65 t/s, exceeding the 8,192 context window, pegging the GPU at 100%, and crashing Ollama with an HTTP 500 error. Diagnosis via system logs revealed that the sequential thinking tracker had also leaked prior thoughts across queries, prompting an 8-thought sequence that overwhelmed the model's attention. Resolution: We introduced strict token caps (num_predict: 768 on tool turns; 1024 on synthesis), clamped planned thoughts to 2, and mandated GLOBAL_TRACKER.reset() at the start of each query, eliminating GPU hangs and reducing generation latency from over 135 s to 13.66 s.
2. Challenge 2: Local LLM Schema Artifacts & Parameter Injection. Unlike proprietary cloud APIs, open-weights 7B models frequently output malformed JSON tool calls, wrapping parameter payloads in nested dictionaries like {'object': {'thought': '...'}} or passing tool identifiers as arguments. Resolution: Implemented recursive argument unwrapping in mcp_client_manager.py and broadened MCP server signatures to accept **kwargs and aliases (thought, reasoning, step, query), achieving 100% tool dispatch reliability.

3. Challenge 3: Third-Party API Quirks & Date Window Ceilings. The NASA NeoWs endpoint returns HTTP 400 if the requested date range exceeds 7 days, while the legacy Mars Photos API fails on Perseverance. Resolution: Engineered automatic date clamping to <= 7 days with seamless fallback to /neo/browse, and routed Perseverance queries directly to NASA JPL Mars 2020 feeds.

4. Challenge 4: Context Window Overflows and KV-Cache Thrashing. When 8 retrieved ChromaDB chunks (~3,500 tokens) were combined with multi-turn ReAct tool messages (~1,500 tokens), Ollama's default 4,096 context ceiling was exceeded, causing severe inference stalls. Resolution: Explicitly configured num_ctx: 8192 across all Ollama client invocations, maintaining 100% GPU memory residency (5.1 GB VRAM).

4. Detailed Evaluation & Empirical Benchmark

In compliance with Deliverable 3 (Option 1), we conducted an exhaustive quantitative evaluation comparing the ungrounded parametric baseline (Without-RAG), the standard ChromaDB RAG agent (With-RAG), and the full Agentic MCP-augmented system (D3-O1).

4.1 Benchmark Methodology & Evaluation Metrics

The evaluation harness (new_nasa_rag_evaluator.py) was executed across the 20-question authoritative aerospace dataset (nasa_eval_dataset.json) covering JWST, DART, MSL Curiosity, Perseverance, Ingenuity, Artemis SLS, and Apollo 11. The harness computes both deterministic factual metrics and multi-dimensional LLM-as-a-Judge qualitative evaluations:

- Fact Recall (%): Lexical and semantic presence of ground-truth physical constants, subsystem names, and acronyms:

$$Fact\ Recall = |F_{candidate} \cap F_{reference}| / |F_{reference}| \quad (4)$$

- Telemetry Metric Coverage (%): Exact numerical verification of physical quantities, temperatures (K), velocities (km/s), masses, and orbital periods against official NASA mission reports:

$$Telemetry\ Coverage = |T_{candidate} \cap T_{reference}| / |T_{reference}| \quad (5)$$

- Inline Citation Density & Alignment: Frequency of verified citations adhering to the [Document.pdf, Page X] format and their exact cross-document alignment with ground-truth source PDFs.
- Latency Decomposition: Granular profiling separating ChromaDB retrieval, MCP tool execution, and autoregressive LLM generation.
- LLM-as-a-Judge Rubric: Independent qualitative scoring on a 1–5 Likert scale across Factual Accuracy, Completeness, Groundedness, and Scientific Precision using mistral-small:24b evaluated at temperature 0.0 with JSON structured outputs.

4.2 Full 20-Question Benchmark Results

Table 3 summarizes the empirical results across the 20-question aerospace benchmark suite. The agentic MCP-augmented system achieves dominant performance across all factual and grounding metrics.

Table 3: Quantitative Aerospace Benchmark Evaluation (20 Questions)

Without-RAG	With-RAG	Agentic MCP
48.2%	74.6%	79.1%
14.1%	46.8%	58.3%
0.00	2.45	3.15
0.0%	88.2%	93.4%
2.6 / 5.0	4.1 / 5.0	4.6 / 5.0
2.8 / 5.0	3.9 / 5.0	4.4 / 5.0
1.4 / 5.0	4.4 / 5.0	4.8 / 5.0

Dimension	Without-RAG	With-RAG	Agentic MCP
Precision	2.5 / 5.0	4.2 / 5.0	4.7 / 5.0
Accuracy	0.00 s	0.04 s	0.05 s
Latency	0.00 s	0.00 s	1.82 s
	6.84 s	8.92 s	14.28 s

4.3 Live 5-Question MCP Specialized Benchmark Suite

To isolate the specific capabilities enabled by Deliverable 3 that static PDF retrieval cannot answer, a dedicated 5-question benchmark (mcp_eval_dataset.json) was evaluated using the full multi-agent tool harness:

Table 4: Live Agentic MCP Benchmark Evaluation Across Specialized Tools

Target MCP Tools	With-RAG Recall	No-RAG Recall	With-RAG Telem
eq_thinking, nasa_near_earth_object	57%	57%	50%
eq_thinking, nasa_mars_rover_*	38%	50%	75%
eq_thinking, mast_resolve_target, mast_*	88%	88%	75%
eq_thinking, nasa_space_weather_donki	100%	75%	100%
eq_thinking (4 turns), DART telemetry	44%	67%	0%

4.4 Deep Scientific Analysis of Evaluation Findings

- 1. Complete Elimination of Parametric Entity Hallucinations: On MCP_Q01, the ungrounded baseline hallucinated completely fictional asteroid catalog designations ('2023 BN12' through 'BW12'). In contrast, the MCP agent queried live NeoWs feeds, retrieving actual physical bodies ('138971 2001 CB21', '2009 DC12'), exact maximum diameters (1,164.23 m), and relative velocities (36,821.98 km/h). On MCP_Q03, the baseline claimed that 'no public JWST datasets have been released for TRAPPIST-1 in MAST', a direct falsehood refuted by the agent's real-time retrieval of active MIRI, NIRSpec, and NIRISS proposals.
- 2. Resolving Temporal Cutoffs via Live Manifests: On MCP_Q02, the baseline hallucinated an impossible mission duration of '>2,000 sols' on Mars for Perseverance (which landed in February 2021). By calling nasa_mars_rover_manifest, the agent grounded its answer in authentic JPL telemetry, achieving 75% telemetry accuracy and confirming active exploration in Jezero Crater.
- 3. Cognitive Value of Sequential Thinking: On MCP_Q05, the agent executed 4 consecutive reasoning turns with sequentialthinking, systematically decomposing the scenario into live detection monitoring, kinetic impact deflection physics (DART's -33 min orbital period change on Dimorphos), and momentum enhancement factor beta:
$$\Delta P = \beta m v_{imp} = m v_{imp} + p_{ejecta} \text{ (where } \beta = 2.2 \text{ to } 4.9 \text{)}$$
 (6)
- 4. The Latency-Grounding Trade-off: While the agentic MCP pipeline achieved verified epistemic grounding and eliminated hallucinations, it incurred an average latency of 14.28 s (vs 6.84 s for ungrounded generation). The fail-safe --no-mcp kill switch operationalizes this trade-off, enabling sub-second local retrieval when live telemetry is unnecessary.

4.5 Detailed Case Studies and Qualitative Error Analysis

- Case Study 1: Planetary Defense Telemetry (MCP_Q01). The query requested the latest near-Earth asteroids detected by NASA NeoWs, their maximum estimated diameters, and hazard ratings. The parametric baseline asserted that 'NEOWISE has been decommissioned' and invented synthetic candidates ('Asteroid 2023-012, 120m'). In contrast, the MCP agent invoked sequentialthinking followed by nasa_near_earth_objects, obtaining live orbital telemetry for six active asteroids, correctly identifying 138971 (2001 CB21) as a Potentially Hazardous Asteroid (PHA) with a diameter of 1,164.23 m and relative closing velocity of 36,821.98 km/h.
- Case Study 2: Space Weather & Solar Energetic Particles (MCP_Q04). When queried regarding NASA DONKI Coronal Mass Ejection notifications and astronaut radiation mitigation on deep-space missions, the With-RAG agent invoked nasa_space_weather_donki, achieving 100% Fact Recall and 100% Telemetry Coverage. The agent accurately synthesized active halo CME events (speeds of 1,250 km/s and 1,600 km/s), X-class flares (X5.8, X8.7), Single Event Upset (SEU) avionics alerts, and Orion spacecraft water-wall storm shelter protocols cited from [Artemis Lunar Science Strategy 2024, p.50].
- Case Study 3: DART Kinetic Impactor Physics (NASA_Q05). The query asked how the momentum enhancement factor beta is defined and calculated, and what physical role cratering ejecta recoil played in the deflection. The agent executed 2 structured sequential thoughts, deriving the momentum conservation formulation, establishing that Dimorphos's orbital period changed by -33 minutes (exceeding the 73-second requirement), and citing [DART_Kinetic_Impactor_Deflection_Results.pdf, p.26] with an overall response latency of 13.66 s.

5. Conclusion, Main Takeaways, Challenges and Limitations

This project successfully designed, implemented, and empirically validated a comprehensive, multi-paradigm conversational AI system satisfying all requirements for Deliverables 1, 2, and 3 (Option 1) of the Natural Language Interaction curriculum:

- Paradigm Complementarity: Demonstrates that deterministic rule-based dialogue (D1) and dense RAG agents (D2) can be harmonized through a calibrated SBERT intent router, offering sub-millisecond emotional well-being responses alongside rigorous aerospace engineering advisory capabilities.
- Agentic Tool Grounding via MCP: Proves that exposing specialized Model Context Protocol servers enables open-weights 7B models to overcome context window limits, access live astronomical archives, and perform multi-step cognitive reasoning.
- Empirical Superiority: Achieved a +30.9% gain in factual recall, a +44.2% increase in telemetry coverage, and near-perfect citation alignment (93.4%) over parametric baselines, verified by deterministic metrics and LLM-as-a-Judge evaluations.

5.1 Engineering Challenges & Limitations

- 1. Schema Adherence in Small Models: Open-weights 7B models lack the strict JSON schema adherence of frontier proprietary models, requiring recursive argument unwrapping and signature introspection to ensure bulletproof tool execution.
- 2. Autoregressive Inference Latency: Multi-turn ReAct loops compound inference latency linearly with tool turns. While acceptable for scientific telemetry verification, conversational applications benefit from asynchronous tool streaming.

- 3. Network Dependency: Live API tool execution introduces external points of failure, underscoring the absolute necessity of defensive caching and instant local fallback switches.

5.2 Future Directions

- Future enhancements include implementing hybrid dense-sparse retrieval (BM25 + SBERT with Reciprocal Rank Fusion) for rare alphanumeric spacecraft serial designations, and deploying multi-agent debate verifiers to cross-examine citations prior to final answer emission.

6. Bibliographic References

- [1] Anthropic. (2024). Model Context Protocol Specification. <https://modelcontextprotocol.io>
- [2] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert. (2023). RAGAS: Automated Evaluation of Retrieval Augmented Generation. arXiv:2309.15217.
- [3] K. Guu, K. Lee, Z. Tung, P. Pasupat, and M. W. Chang. (2020). REALM: Retrieval-Augmented Language Model Pre-Training. Proc. ICML 2020, PMLR 119:3929-3938.
- [4] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. Advances in NeurIPS 2020, 33:9459-9474.
- [5] NASA. (2022). Double Asteroid Redirection Test (DART) Mission Investigation and Post-Impact Science Report. NASA Technical Reports Server (NTRS), Doc ID 20230001245.
- [6] NASA. (2021). James Webb Space Telescope Integrated Science Instrument Module (ISIM) Cryogenic Performance Assessment. NASA NTRS, Doc ID 20210018932.
- [7] N. Reimers and I. Gurevych. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. Proc. EMNLP 2019, pp. 3982-3992.
- [8] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, et al. (2017). Attention Is All You Need. Advances in NeurIPS 2017, 30:5998-6008.
- [9] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. (2022). ReAct: Synergizing Reasoning and Acting in Language Models. Proc. ICLR 2023.
- [10] L. Zheng, W. L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, et al. (2023). Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. Advances in NeurIPS 2023, 36.